

PyTorch Introduction

Deep Learning



- Torch was introduced in 2002 – an open-source machine learning library and scientific computing framework based on Lua.
- Developed at the Idiap Research Institute at EPFL in Lausanne, Switzerland.
- The main developers were Ronan Collobert, Samy Bengio and Johnny Mariétho.
- In 2016 the front-end was converted by Soumith Chintala, Adam Paszke, Sam Gross and Gregory Chanan at Facebook (now Meta) from Lua to Python, and renamed PyTorch.



Model training steps (same as TensorFlow)

- Load the data.
- Construct the network.
- Train the network.
- Analyze the results.



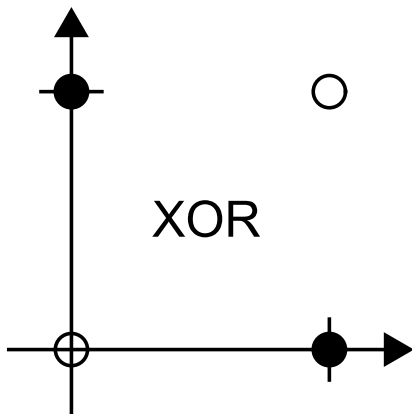
- PyTorch has the same formats for tensors as TensorFlow (features, samples, timesteps and channels).
- PyTorch tensors can be created directly from data, using the `torch.tensor()` command.
- Can be created from a numpy array, using the `torch.from_numpy()`.



- Each element of a network input is called a **feature**.
- The data set will consist of Q **samples** of inputs.
- If each input is a vector (tabular) with R features, then the data set is a (Q, R) , or (samples, features), NumPy tensor.
- If the input is a time series, the input is a 3D tensor of the form (samples, timesteps, features).
- For images, the 4D input tensor form is (samples, height, width, channels), where channels are usually colors.
- For videos, the 5D input tensor form is (samples, timesteps, height, width, channels).



XOR test problem



Generate XOR data

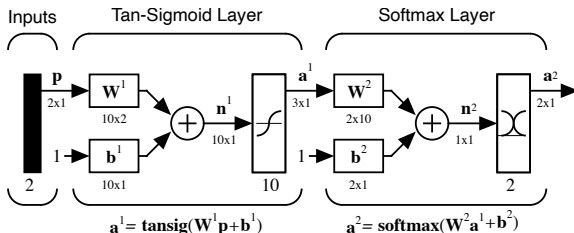
```
import numpy as np
import torch

p = torch.tensor([[0, 0], [0, 1], [1, 0],
    ↪ [1, 1]], dtype=torch.float32)
t = torch.tensor([0, 1, 1, 0])
print(p)
print(t)
```

```
tensor([[0., 0.],
        [0., 1.],
        [1., 0.],
        [1., 1.]])
tensor([0, 1, 1, 0])
```



Constructing the model with Sequential



```
model = torch.nn.Sequential(  
    torch.nn.Linear(2, 10),  
    torch.nn.Tanh(),  
    torch.nn.Linear(10, 2))  
print(model)
```

```
Sequential(  
  (0): Linear(in_features=2, out_features=10, bias=  
    ↪ True)  
  (1): Tanh()  
  (2): Linear(in_features=10, out_features=2, bias=  
    ↪ True))
```



Constructing the model using the model subclass method

```
class TwoLayer(torch.nn.Module):
    def __init__(self):
        super().__init__()
        self.dense1 = torch.nn.Linear(2, 10)
        self.tanh = torch.nn.Tanh()
        self.dense2 = torch.nn.Linear(10, 2)
    def forward(self, x):
        x = self.dense1(x)
        x = self.tanh(x)
        x = self.dense2(x)
        return x
model = TwoLayer()
print(model)
```

```
TwoLayer(
  (dense1): Linear(in_features=2, out_features=10,
    ↪ bias=True)
  (tanh): Tanh()
  (dense2): Linear(in_features=10, out_features=2,
    ↪ bias=True))
```



- We need to select the optimizer.
- First argument of the optimizer contains the variables to be adjusted.
- You can also specify optimizer options, such as learning rate.

```
optimizer = torch.optim.Adam(model.  
    ↪ parameters(), lr=0.001)
```

- You also need to define the loss function.

```
loss_fn = torch.nn.CrossEntropyLoss()
```



- 1 Zero the gradient
- 2 Make a forward pass through the network
- 3 Calculate the loss
- 4 Compute the gradient
- 5 Update the weights



Training the network

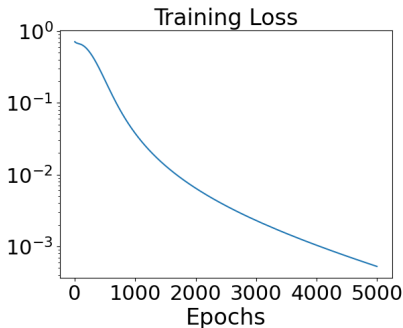
```
#Check if a GPU is available, and move model and data  
device = 'cuda' if torch.cuda.is_available() else '  
    ↪ cpu'  
model.to(device)  
p = p.to(device)  
t = t.to(device)  
epochs = 5000  
loss_values = []  
for epoch in range(epochs):  
    # Zero the gradient  
    optimizer.zero_grad()  
    # Forward pass through the network  
    output = model(p)  
    # Calculate the loss  
    loss = loss_fn(output, t)  
    # Compute the gradient  
    loss.backward()  
    # Save the loss values for plotting  
    loss_values.append(loss.cpu().detach().numpy())  
    # Update the weights  
    optimizer.step()
```



Convergence plot

We saved the loss values, so that we could view the process.

```
import matplotlib.pyplot as plt
plt.semilogy(range(epochs), loss_values)
plt.title('Training Loss')
plt.xlabel('Epochs')
plt.show()
```



```
n = model(p)
softmax = torch.nn.Softmax(dim=1)
a = softmax(n)
print(a)
```

```
tensor([[9.9967e-01,  3.2964e-04],
        [5.4206e-04,  9.9946e-01],
        [2.5702e-04,  9.9974e-01],
        [9.9960e-01,  4.0006e-04]], device=
        ↪ 'cuda:0', grad_fn=<
        ↪ SoftmaxBackward0>)
```

We added softmax, because it was not included in the model.



- Before network training comes Extract, Transform and Load (ETL).
- First, the data are taken from one or multiple files, which may be distributed across multiple machines.
- Next, the data is transformed (normalizing, augmenting by rotating or scaling images, adding noise, etc.)
- Finally, the data is loaded into the training process, often in minibatches.
- `torch.utils.data.Dataset` defines what the data are and how to access them.
- `torch.utils.data.DataLoader` defines how to load data efficiently for training or inference.



- 1 Represents the data and enables access to individual samples.
- 2 Each item of the dataset has an input and a target.
- 3 Defines how to load and preprocess individual data points.
- 4 Implements `__getitem__()` to retrieve a single sample and `__len__()` for the total number of samples.
- 5 Focuses on data representation and access.



- ① Wraps a Dataset and provides utilities for batching, shuffling, and parallel data loading
- ② Handles the iteration over the dataset, creating batches, and optionally shuffling
- ③ Allows easy specification of batch size, number of worker processes, etc.
- ④ Focuses on efficiently feeding data to the model during training/evaluation



```
import pandas as pd  
sample_df = pd.read_csv('SampleDF.csv')
```

Extract two columns that we will use as inputs and targets.

```
P = np.array(sample_df['FVC'])  
T = np.array(sample_df['Percent'])
```



Define a Dataset class

```
class SimpleDataset(torch.utils.data.Dataset):  
    def __init__(self, P, T):  
        # convert into PyTorch tensors  
        self.P = torch.tensor(P, dtype=torch.float32)  
        self.T = torch.tensor(T, dtype=torch.float32)  
  
    def __len__(self):  
        # this should return the size of the dataset  
        return len(self.P)  
  
    def __getitem__(self, idx):  
        # return one input and one target sample from  
        ↪ the dataset  
        features = self.P[idx]  
        target = self.T[idx]  
        return features, target
```



Load and iterate the Dataset

```
dataset = SimpleDataset(P, T)
```

Wrap the data in a DataLoader.

```
loader = torch.utils.data.DataLoader(dataset, shuffle  
    ↪ =True, batch_size=4)
```

Iterate the DataLoader.

```
for feat, targ in loader:  
    print('Features: {}, Target: {}'.format(feat, targ  
        ↪ ))  
    break
```

```
Features: tensor([2603., 2762., 1537.,  
    ↪ 4738.]), Target: tensor([ 71.1591,  
    ↪ 72.4174, 63.5281, 113.0464])
```

